

Terence Zhi Liu

Machine Learning / Platform Engineering

[linkedin.com/in/terence-z-liu](https://www.linkedin.com/in/terence-z-liu)

terence.z.liu@gmail.com

(419) 936-4932

For detailed projects and write-ups, please visit terencezl.github.io and terencezl.github.io/blog.

EXPERIENCE

- **Head of Machine Learning / Principal Engineer | Clearview AI, NY** 2021 - Present
 - Created a facial recognition [algorithm](#) that ranked **US #1** and **World #2** by [NIST](#).
 - Researched & implemented a **deca-B** vector database **100+ times** the SOTA throughput. Fused high-profile open-source C++ libraries to push the industry boundary. Near-term 1/10 server cost & much higher capacity. Long-term \$multi-M annual savings.
 - Perfected efficient model training using tens of nodes with machine & job orchestration tools.
 - Drove model efficiency using distillation training, quantization, pruning & trimming, and hardware acceleration. Streamlined deployment with model encryption and serving engines.
 - Led cross-functional collaboration to design & review multiple compact ML data pipelines & inference infra for images/videos, resulting in \$100k cost reduction per recurring batch job.
 - Established foundational ML practices, tooling & platform and built a team of ML engineers.
 - Planned & executed product development & GTM of [PAD](#) feature with team.
 - Co-wrote marketing materials and public letters. Engaged with policy officials.
 - *Liu, Ton-That, Scalable Training Data Pipeline... (2022) U.S. Patent No. 11,443,553*
 - **New:** PyTorch, OpenCV, ONNX Runtime, OpenVINO, TensorRT, AWS SageMaker, GCP Vertex AI, [Ray](#), Airflow, MLFlow, DVC, [Faiss](#), [hnswnlib](#), FastAPI, Socket.IO, Protobuf & gRPC, [RocksDB](#), [Datadog](#), [GitLab CI/CD](#)
- **Senior Software Engineer, Platform | Bloomberg LP, NY** 2017 – 2021
 - Maintained & monitored widely-used bare-metal SFTP infrastructure with 7M daily logins.
 - Designed & implemented reliable account management, auth, routing, caching, messaging for cloud-based next-gen SFTP. Wrote OS-level modules and web servers for auth interfacing with OpenSSH. Managed deployment on internal Kubernetes-based PaaS.
 - Leveraged S3 Storage as SFTP subsystem with file system emulation and cross-data-center replication & failover.
 - Spearheaded successful multi-year high-stake account migration as technical lead.
 - **New:** C++, Golang, OpenSSH, Flask, SQLAlchemy, Postgres, Redis, MongoDB, RabbitMQ, Kafka, Flink, Spark, Grafana, Splunk, Humio, Jenkins, Chef, Docker, Kubernetes, OpenStack
- **Researcher & Developer | University of Toledo, OH** 2013 – 2017
 - Specialized in materials simulations with parallel computing clusters. 14 [papers](#) published.
 - Calculated electronic ground states with gradient descent & residual minimization schemes, (non-)linear regression - similar routines as in ML/DL frameworks. Had various ML [projects](#).
 - Authored open-source projects: [ScriptsForVASP](#), [pydass_vasp](#), [pyvasp-workflow](#).
 - Built a materials database [website](#) with a modern stack, supporting tabulation/graphing, user auth/contribution. Project helped get a \$100k research grant.
 - *Linux, Python, R, Java, NumPy, SciPy, scikit-learn, XGBoost, LightGBM, pandas, matplotlib*

EDUCATION

- **Ph.D. in Computational Physics | University of Toledo, OH | GPA 4.0/4.0** 2012 – 2017
[Dissertation](#): A First-Principles Simulation Approach with Cluster Expansion Predictive Models
- **B.S. in Materials Physics | Nanjing University, China** 2008 – 2012
Trained in computer science, statistics, operations research, math, physics, etc.